



The Abdus Salam  
International Centre  
for Theoretical Physics



# Use of QE in HPC: overview of implementation and usage of the QE-GPU

**Ivan Girotto – [igirotto@ictp.it](mailto:igirotto@ictp.it)**

Information & Communication Technology Section (ICTS)

International Centre for Theoretical Physics (ICTP)

**... in collaboration with Filippo Spiga (Cambridge University)**



# OUTLINE

- Technological Background
- Development foundation
- Usage
- Performance Analysis
- Conclusions



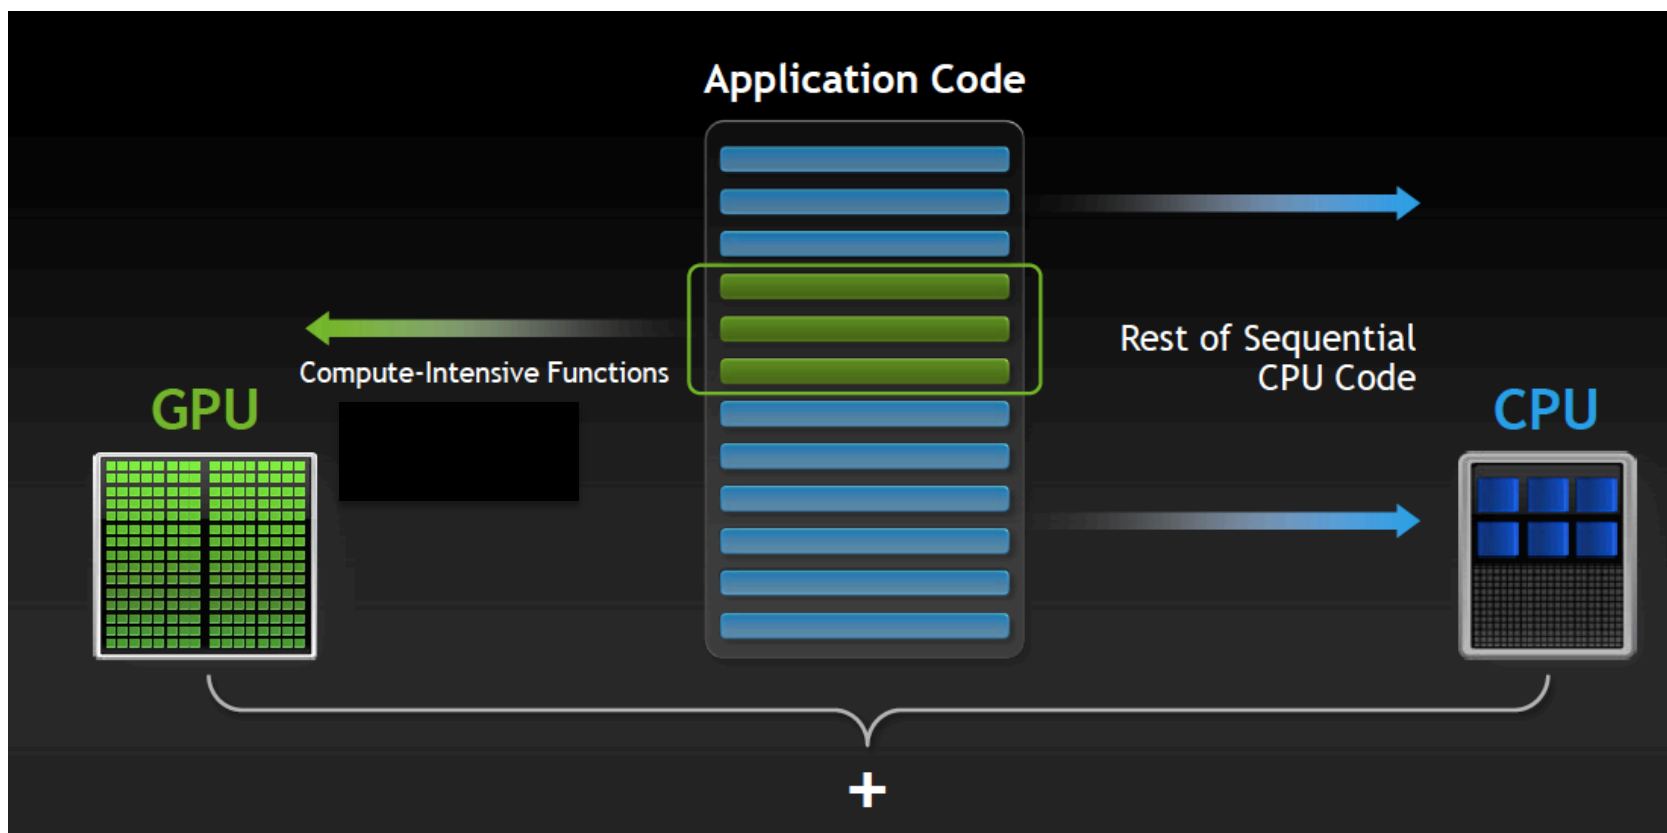
The Abdus Salam  
**International Centre  
for Theoretical Physics**

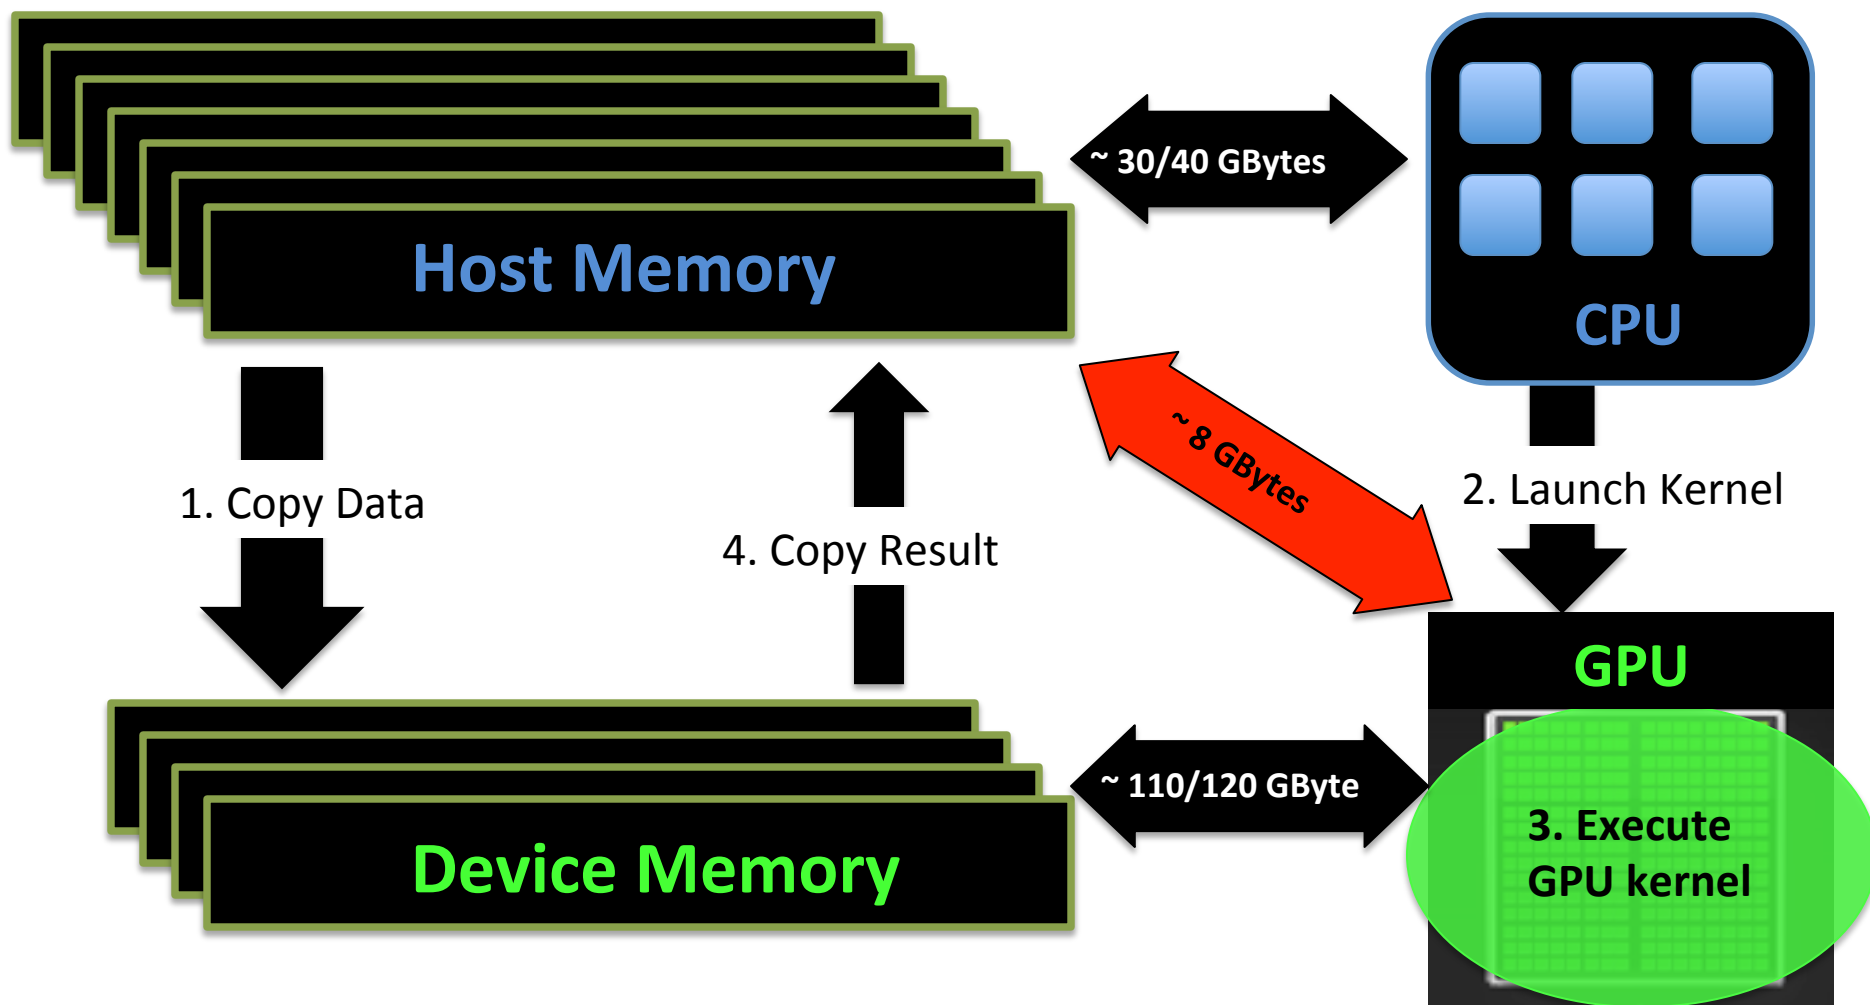


**IAEA**  
International Atomic Energy Agency

# TECHNOLOGICAL BACKGROUND

# The General Concept of Accelerated Computing







# Why Does GPU Accelerate Computing?

- Highly scalable design
- Higher aggregate memory bandwidth
- Huge number of low frequency cores
- Higher aggregate computational power
- Massively parallel processors for data processing



# Why Does GPU Not Accelerate Computing?

- PCI Bus bottleneck
- Synchronization weakness
- Extremely slow serialized execution
- High complexity
  - SPMD(T) + SIMD & Memory Model
- People forget about the Amdahl's law
  - accelerating only the 50% of the original code, the expected speedup can get at most a value of 2!!



# Higher aggregate computational power

- Do we really ... need it? ... have it available?
- Can we really exploit it?
- The DP peak of performance is done as follows:
  - #operations per clock cycle x frequency x #cores
  - we automatically reduce the DP power if we partially use the accelerator
- How much is my GPU better than my CPU?
- A relevant outcome is always a good balance of those issues





The Abdus Salam  
**International Centre  
for Theoretical Physics**



**IAEA**  
International Atomic Energy Agency

# ENABLING OF PWSCF ON GPU



# The Grounds

- The **almost** “perfect” storm
- PRACE-1IP and the Irish Experience
- Visibility
  - having a business compliant product helps (grants, funding, PR 😊)
- Targets:
  - enabling of QE to high-end platforms
  - provide an additional feature to the package
  - make it comparable with other codes enabled on GPU platforms

# The Experienced Problems

- Low DP for consumer computing
- Huge effort for competitiveness and maintenance
- No portability even across generation of NVIDIA GPUs!!
- Software integration and validation



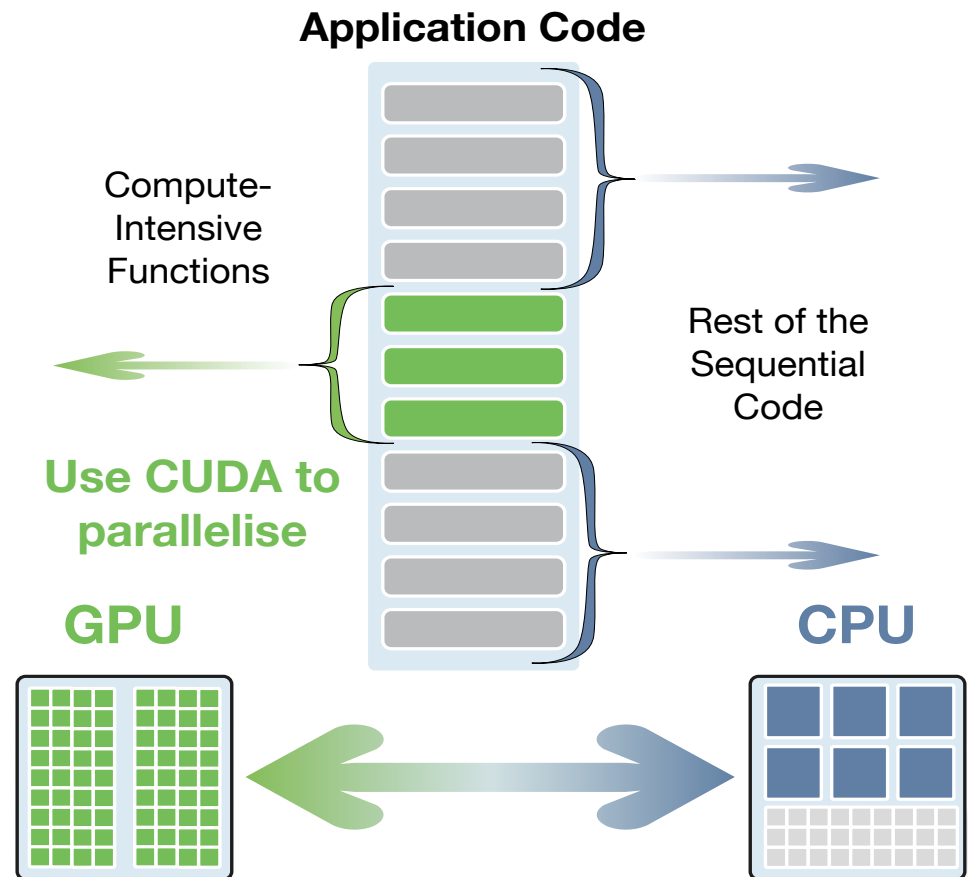
# The QE-GPU repository

- The QE-GPU is now a plug-in maintained in a separated repository
- Please, visit <https://github.com/fspiga/QE-GPU>
- It must be downloaded a part and added into the QE \$ROOT\_DIRECTORY
- We always look for volunteers!!

# Speed-up Scientific Codes

- Directives (OpenACC)
- Libraries
- CUDA (or OpenCL)

**3 Way  
to Accelerate on GPU**



# Development strategy in a Nutshell

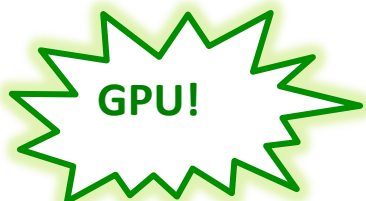
- Important compromise between performances and code re-factoring
- Scalability & Reliability are clear objectives (No trivial Init Module):
  - Multiple processes mapping on single GPU
  - Memory control and management
  - Data transfer overlapping (Pinned memory or not Pinned memory?)
  - Result consistency with the CPU version
- Massive focus on best exploitation of GPU libraries
- Few small routines developed in CUDA (< few hundreds lines)
- Extension to other codes of the QE distribution

# Computational Bottlenecks in QE

- Calculation of density,  $n(\mathbf{r}) = \sum |\psi(\mathbf{r})|^2$  (+ augmentation terms for USPP): FFT + linear algebra (matrix-matrix multiplication)
- Calculation of potential,  $V(\mathbf{r}) = V_{xc}[n(r)] + V_H[n(\mathbf{r})]$ : FFT + operations on real-space grid
- Iterative diagonalization (SCF) / electronic force (CP) calculation,  $H\psi$  products: FFT + linear algebra (matrix-matrix multiplication)
- Subspace diagonalization (SCF) / Iterative orthonormalization of Kohn-Sham states (CP): diagonalization of  $N_e \times N_e$  matrices + matrix-matrix multiplication

\*courtesy of Prof. Paolo Giannozzi

# Levels of parallelism in QE



Images

- Only for Nudged Elastic Band (NEB) calculations

K-points

- Distribution over k-points (if more than one)
- Scalability over k-points (if more than one)
- No memory scaling

Plane-waves

- Distribution of wave-function coefficients
- Distribution of real-grid points
- Good memory scale, good overall scalability, LB

Linear algebra &  
task groups

- Iterative diagonalization (fully-parallel or serial)
- Smart grouping of 3DFFT to reduce *compulsory* MPI communications

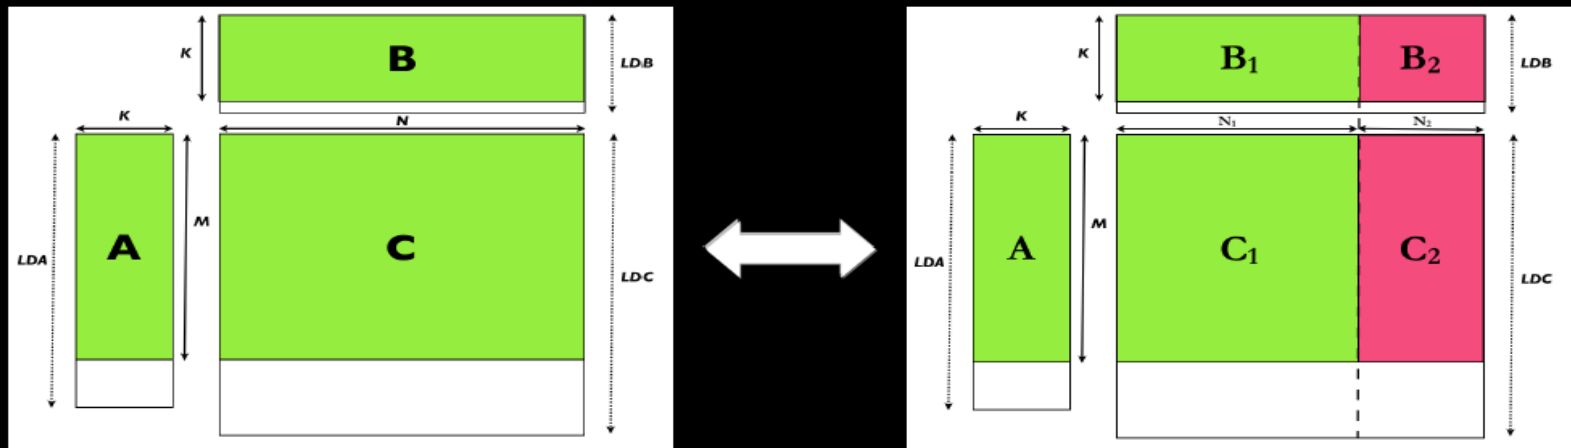
Multi-threaded  
kernels

- OpenMP handled *explicitly* or *implicitly*
- Extend the scaling on multi-core machines with “limited” memory



# Existing work, Dgemm for Linpack HPL

## DGEMM: $C = \alpha A B + \beta C$



$$\text{DGEMM}(A, B, C) = \text{DGEMM}(A, B_1, C_1) \cup \text{DGEMM}(A, B_2, C_2)$$

The idea can be extended to multi-GPU configuration and to handle huge matrices

Find the optimal split, knowing the relative performances of the GPU and CPU cores on DGEMM

Phillips & Fatica [http://www.nvidia.com/content/GTC-2010/pdfs/2057\\_GTC2010.pdf](http://www.nvidia.com/content/GTC-2010/pdfs/2057_GTC2010.pdf)

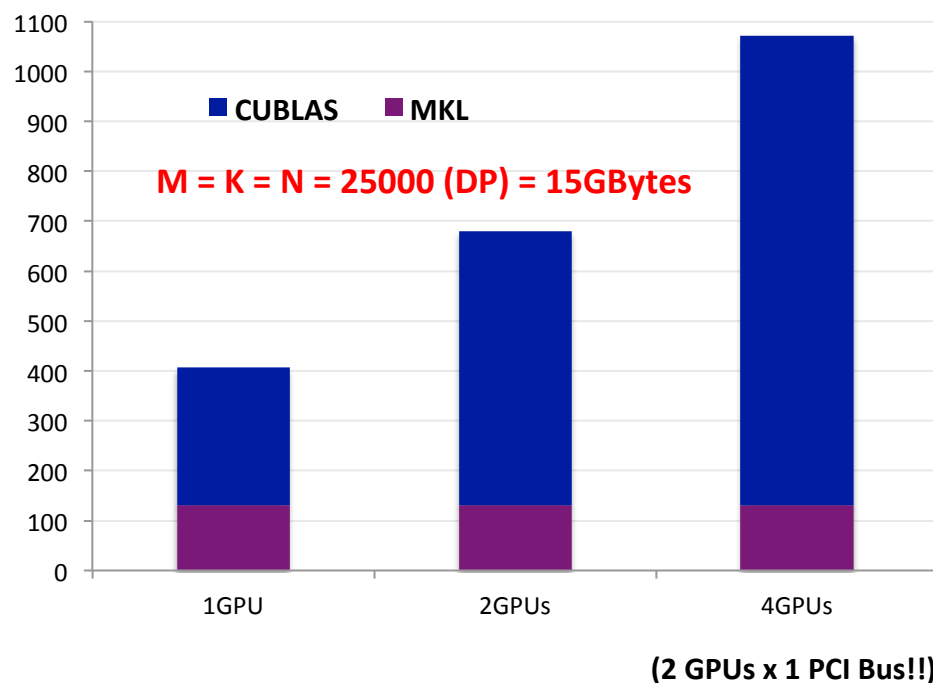
# The starting point: $\Phi$ GEMM

- [\*]GEMM implemented
- Transparent data transferring
- Recursive splitting for big matrixes (no limited by GPU Memory)
- Special-K strategy for rectangular matrixes
- Almost transparently linkable
- Possible profiling report for each [\*]GEMM call
- Possible mappings CPU-processes:GPUs
  - 1:N => MultiGPU version
  - N:1 => 1<sup>+</sup> MPI processes x single GPU

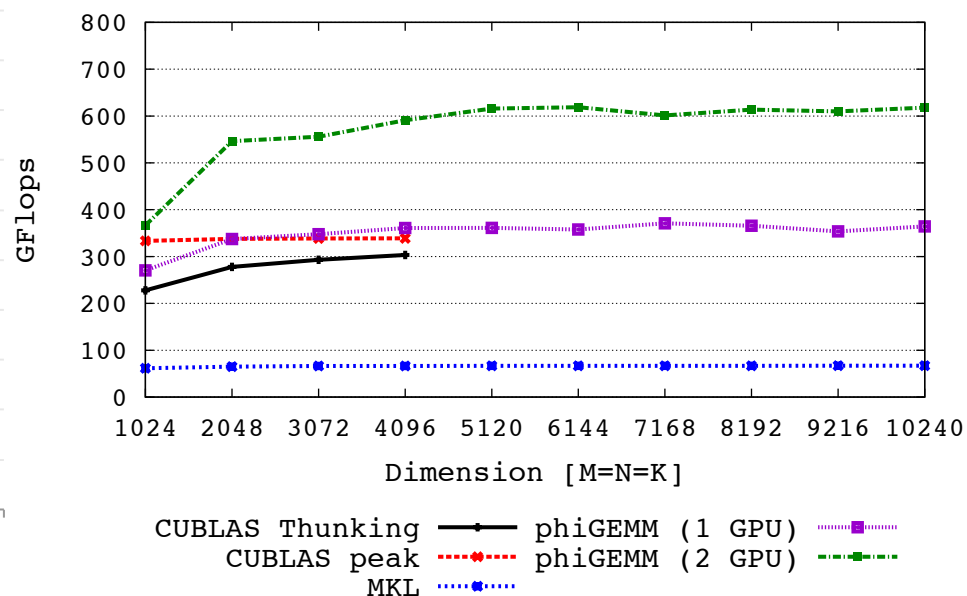


<http://qe-forge.org/gf/project/phigemmm/>

# $\Phi$ GEMM



phiGEMM single- and multi-GPU performance





# CUFFT

- NVIDIA library to interface call to FFTW
- Used only on serial version and if using `-D__USE_3D__FFT` (non distributed plane-waves)
- Good speedup Vs CPU for small number of MPI processes
- C code for wrapping CUFFT completely encapsulated in separated files
- 3DFFT on distributed data still under study: no good efficient solutions found so far

<http://developer.nvidia.com/cuda/cufft>

```

IF ( gamma_only ) THEN
    #if (defined(__CUDA) && ... )
        ierr = vloc_psi_cuda ( lda, dffts%nnr, dffts%nr1x, dffts%nr2x, dffts%nr3x, &
            n, m, psi, vrs(1,current_spin), hpsi, igk(1:), nls(1:), &
            nls(1:), ngms, ngm)
        !
    #else
        CALL vloc_psi_gamma ( lda, n, m, psi, vrs(1,current_spin), hpsi )
    #endif
ELSE IF ( noncolin ) THEN
    !
    CALL vloc_psi_nc ( lda, n, m, psi, vrs, hpsi )
ELSE
    #if (defined(__CUDA) && ... )
        ierr = vloc_psi_cuda_k ( lda, dffts%nnr, dffts%nr1x, dffts%nr2x, dffts%nr3x, &
            n, m, psi, vrs(1,current_spin), hpsi, igk(1:), nls(1:), ngms)
    #else
        CALL vloc_psi_k ( lda, n, m, psi, vrs(1,current_spin), hpsi )
    #endif
END IF

```

# MAGMA

- LAPACK library for GPU systems

```
#if defined(__CUDA) && defined(__MAGMA)
    CALL start_clock( 'MAGMA_DSYGVD' )
    CALL magmaf_dsygvd( 1, 'V', 'U', n, v, ldh, s, &
        ldh, e, work, lwork, iwork, liwork, info )
    CALL stop_clock( 'MAGMA_DSYGVD' )
#else
    CALL start_clock( 'DSYGV' )
    CALL DSYGV( 1, 'V', 'U', n, v, ldh, s, ldh, e, &
        work, lwork, info )
    CALL stop_clock( 'DSYGV' )
#endif
```

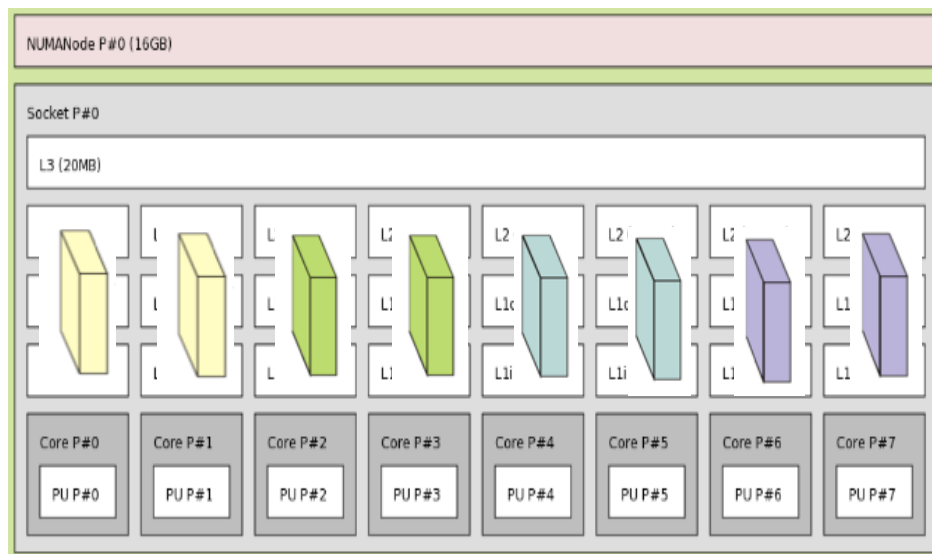
<http://icl.cs.utk.edu/magma/>



## How to use it ...

- The partial porting requires to exploit the overall platform: best performances on CPU and GPU
- All complicated balancing issues presented this morning becomes much relevant for the QE-GPU version
- For instance...



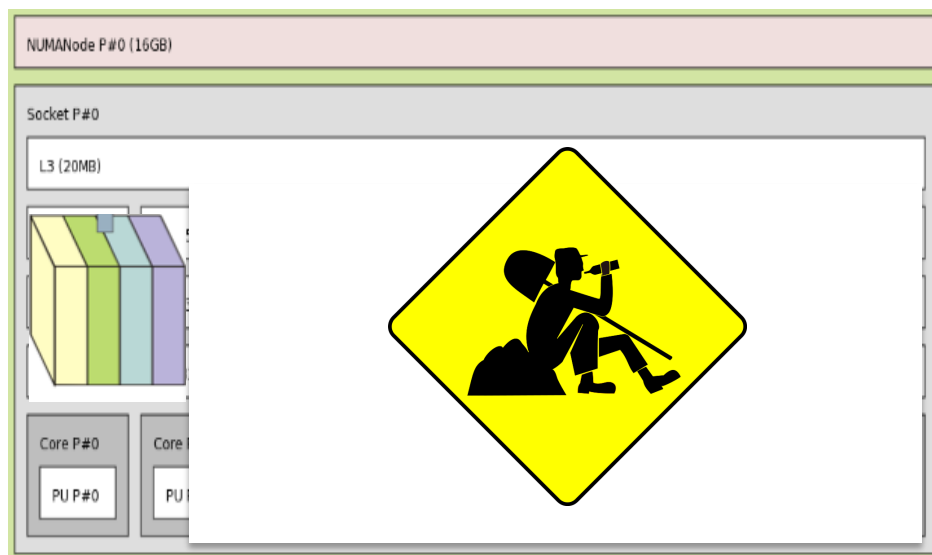


The Intel Xeon E5-2665  
Sandy Bridge-EP 2.4GHz



```
mpirun -np 8 pw-gpu.x -inp input file
```

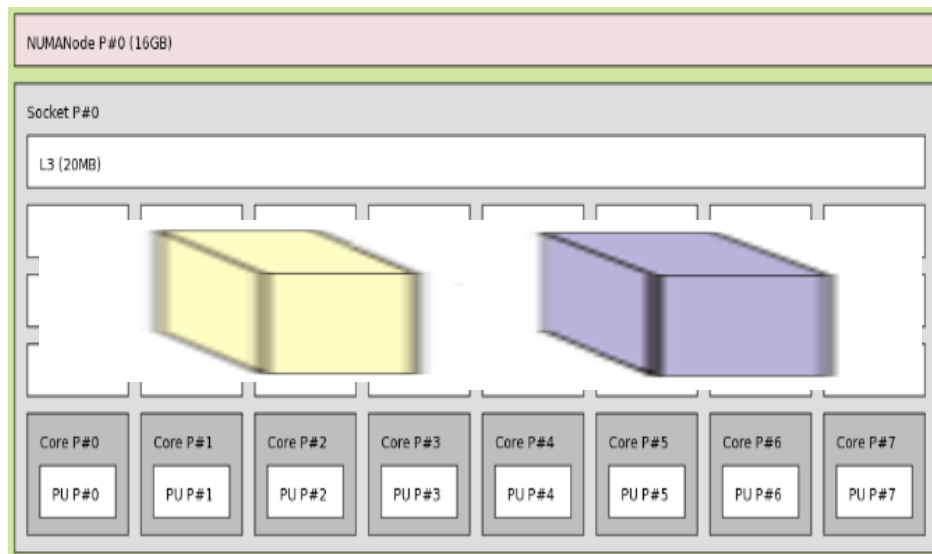




The Intel Xeon E5-2665  
Sandy Bridge-EP 2.4GHz



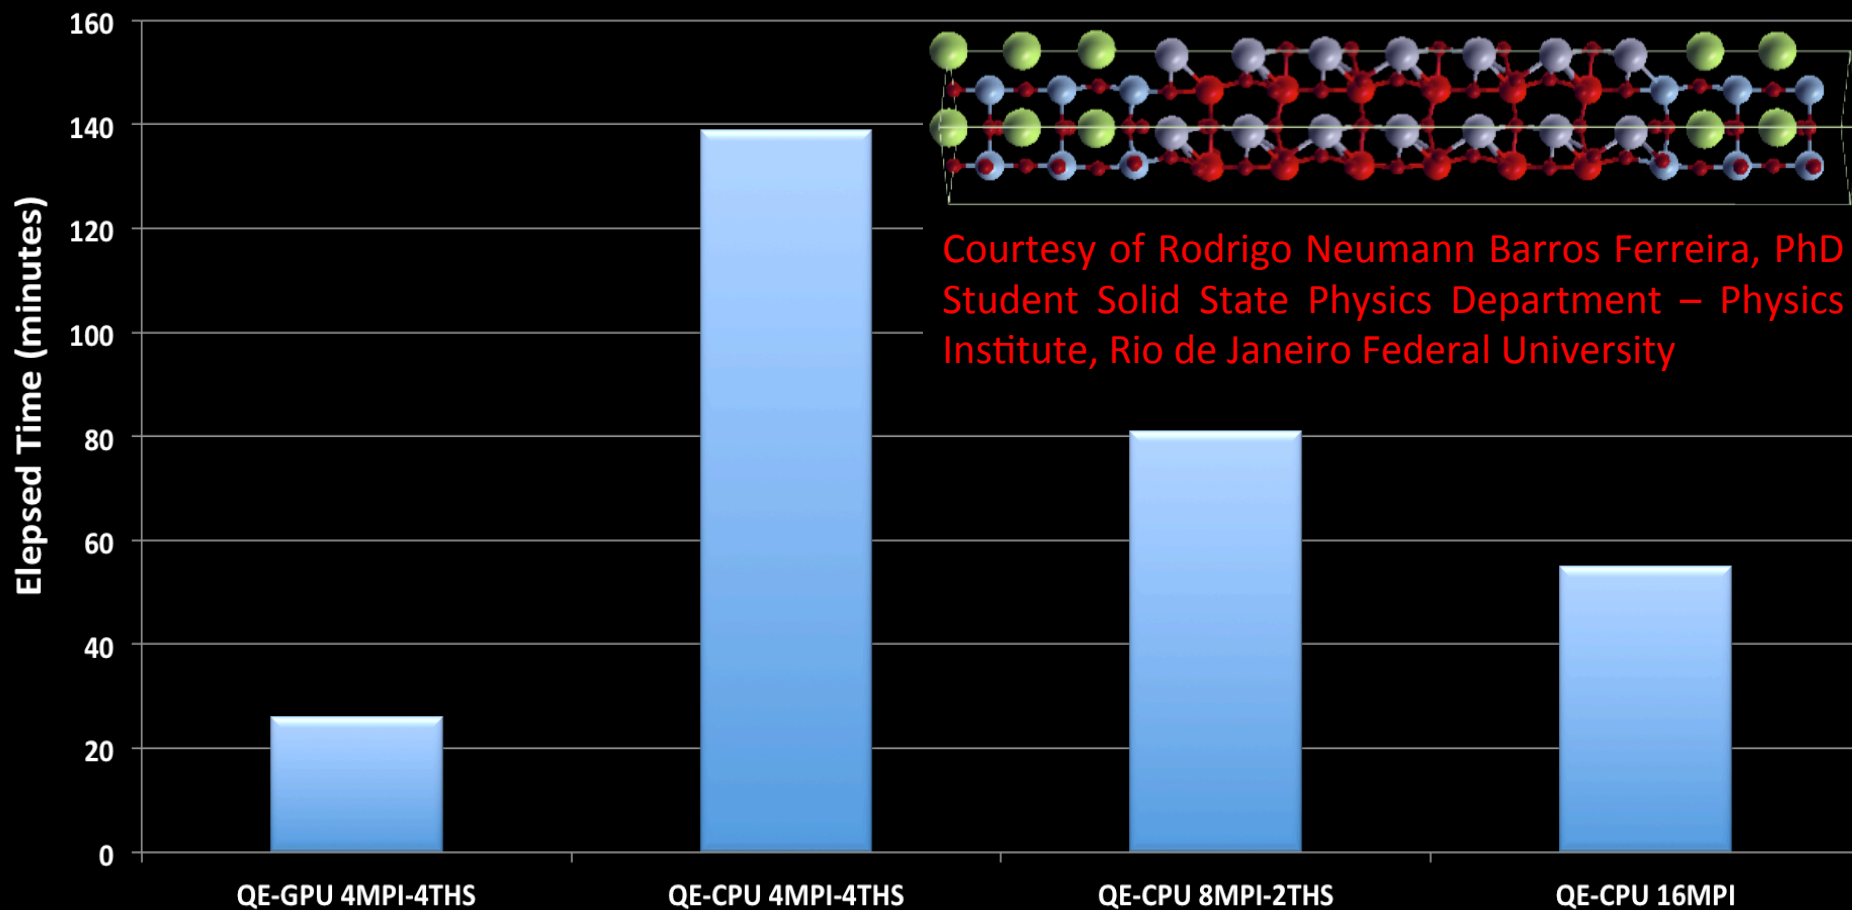
```
mpirun -np 1 pw-gpu.x -inp input file
```

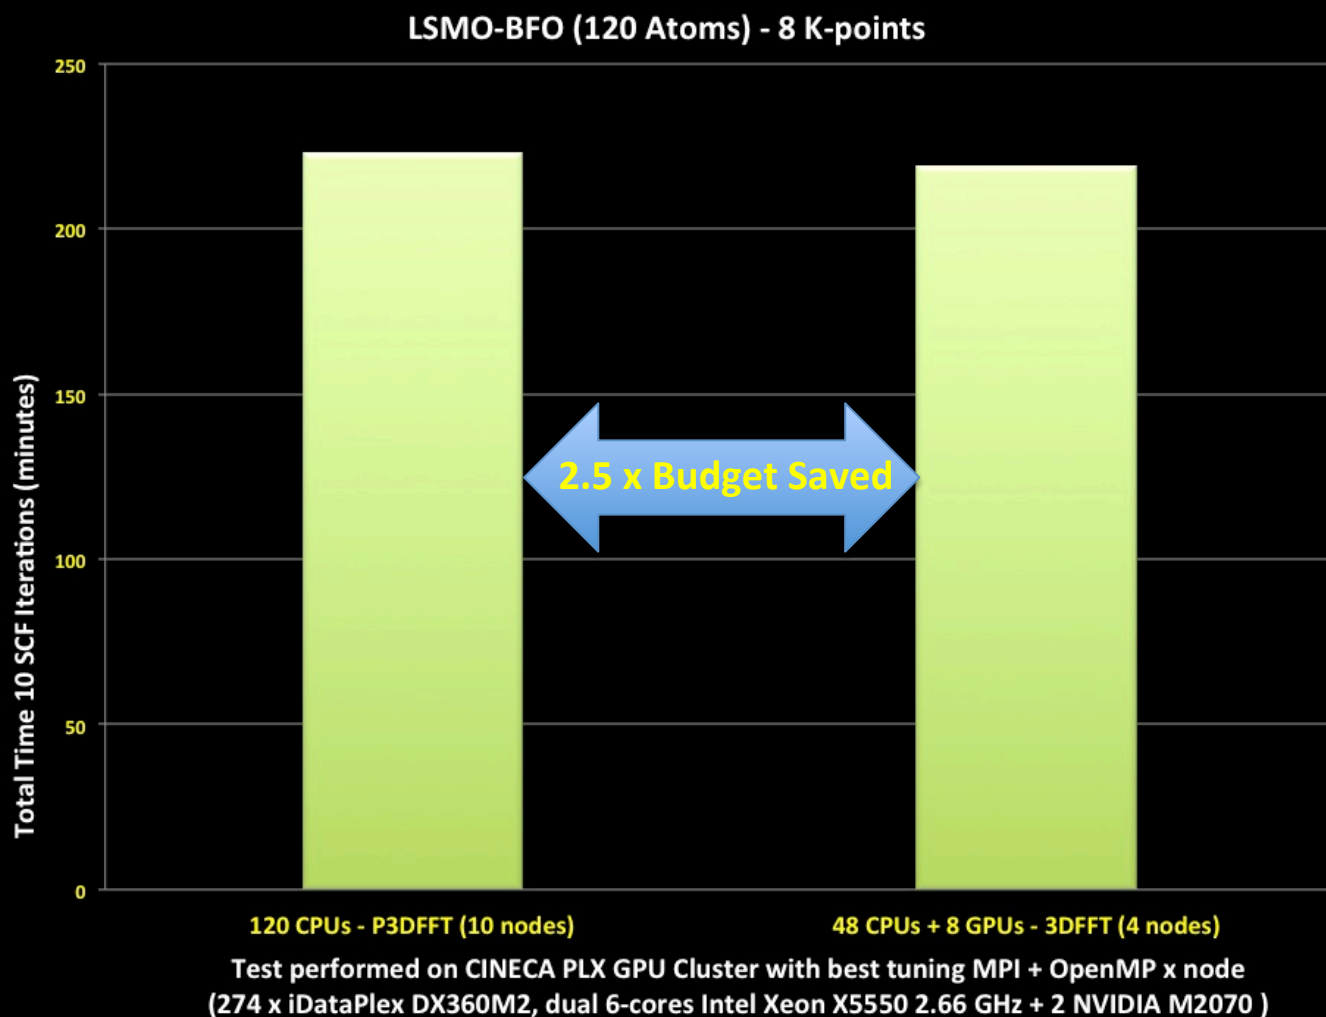


The Intel Xeon E5-2665  
Sandy Bridge-EP 2.4GHz

```
export OMP_NUM_THREADS=4
export OPENBLAS_NUM_THREADS=$OMP_NUM_THREADS
mpirun -np 2 pw-gpu.x -inp input file
```

## scf calculation (few iterations) on Hydra: LSMO BiFeO<sub>3</sub> (BFO)

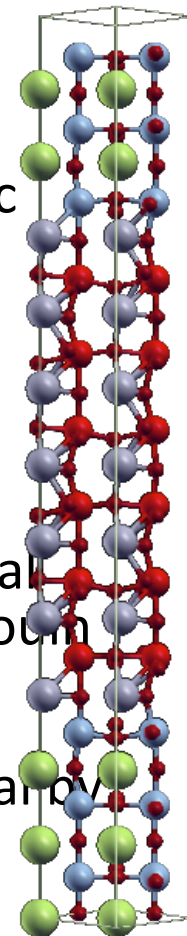


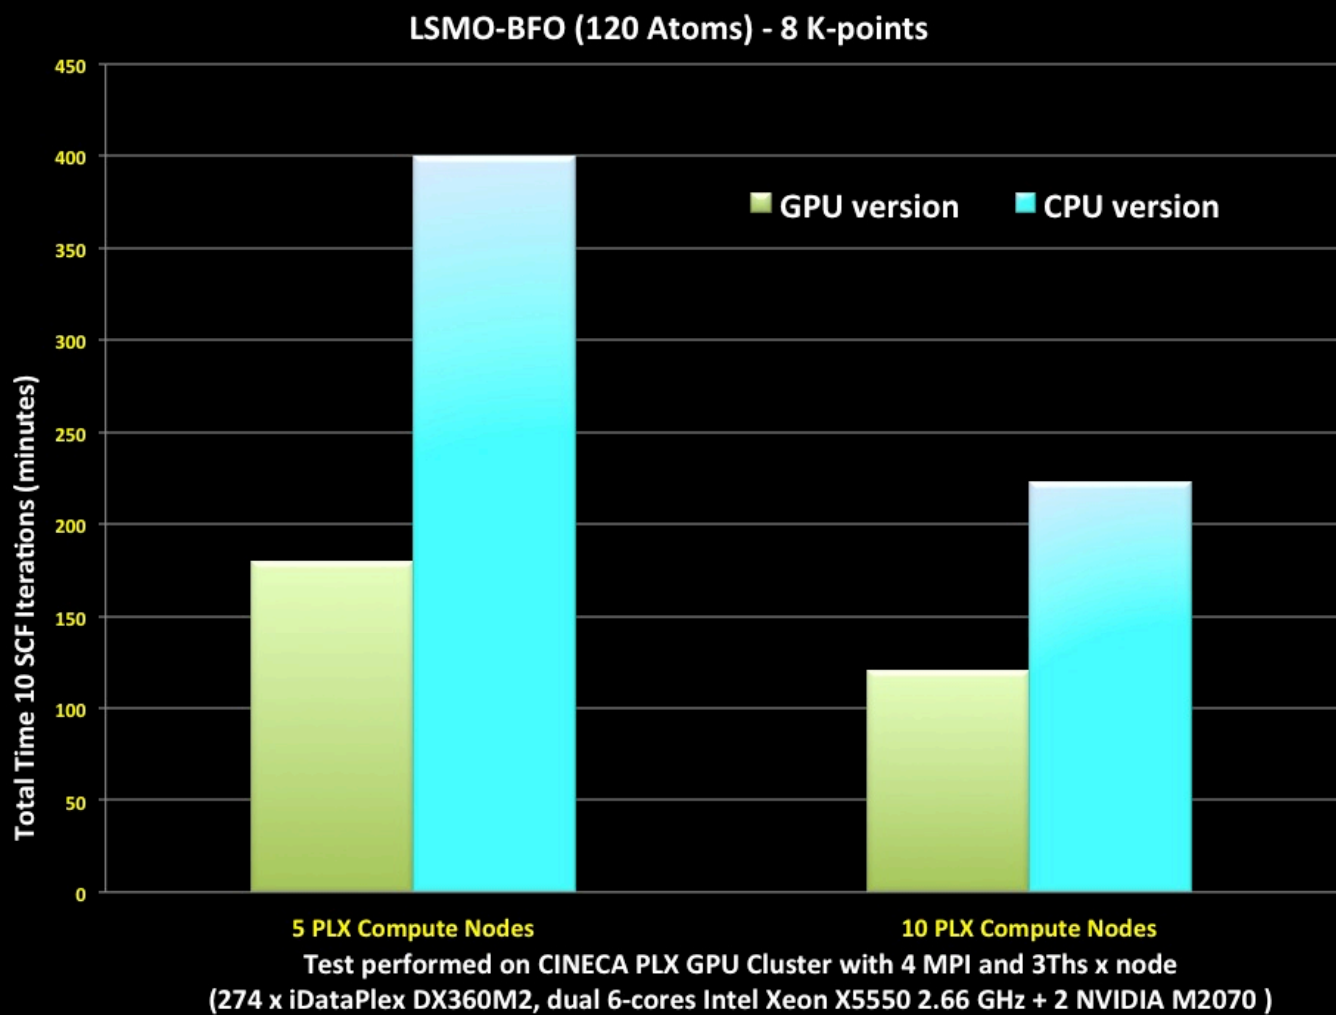


## Performance Results /1

# Best Practice /1

- Scientific case:  $\text{La}_{2/3}\text{Sr}_{1/3}\text{MnO}_3$  /  $\text{BiFeO}_3$  magnetic heterostructure (120 atoms)
- PI: Rodrigo Neumann Barros Ferreira, PhD Student Solid State Physics Department – Physics Institute, Rio de Janeiro Federal University
- Description: 1024 electrons, 615 different quantum-mechanical states considered, **8 k-points** for the integration over the Brillouin zone.
- Goal: exploit QE *pool* parallelism and GPU → keep the FFT local by using `npool=npocs`

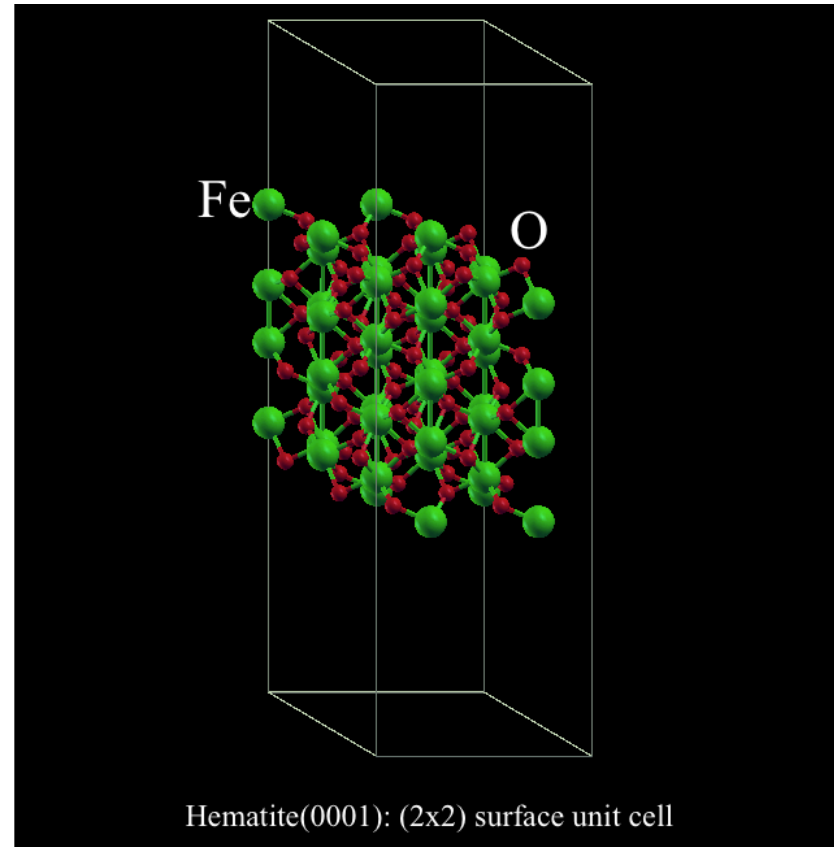




## Performance Results /2

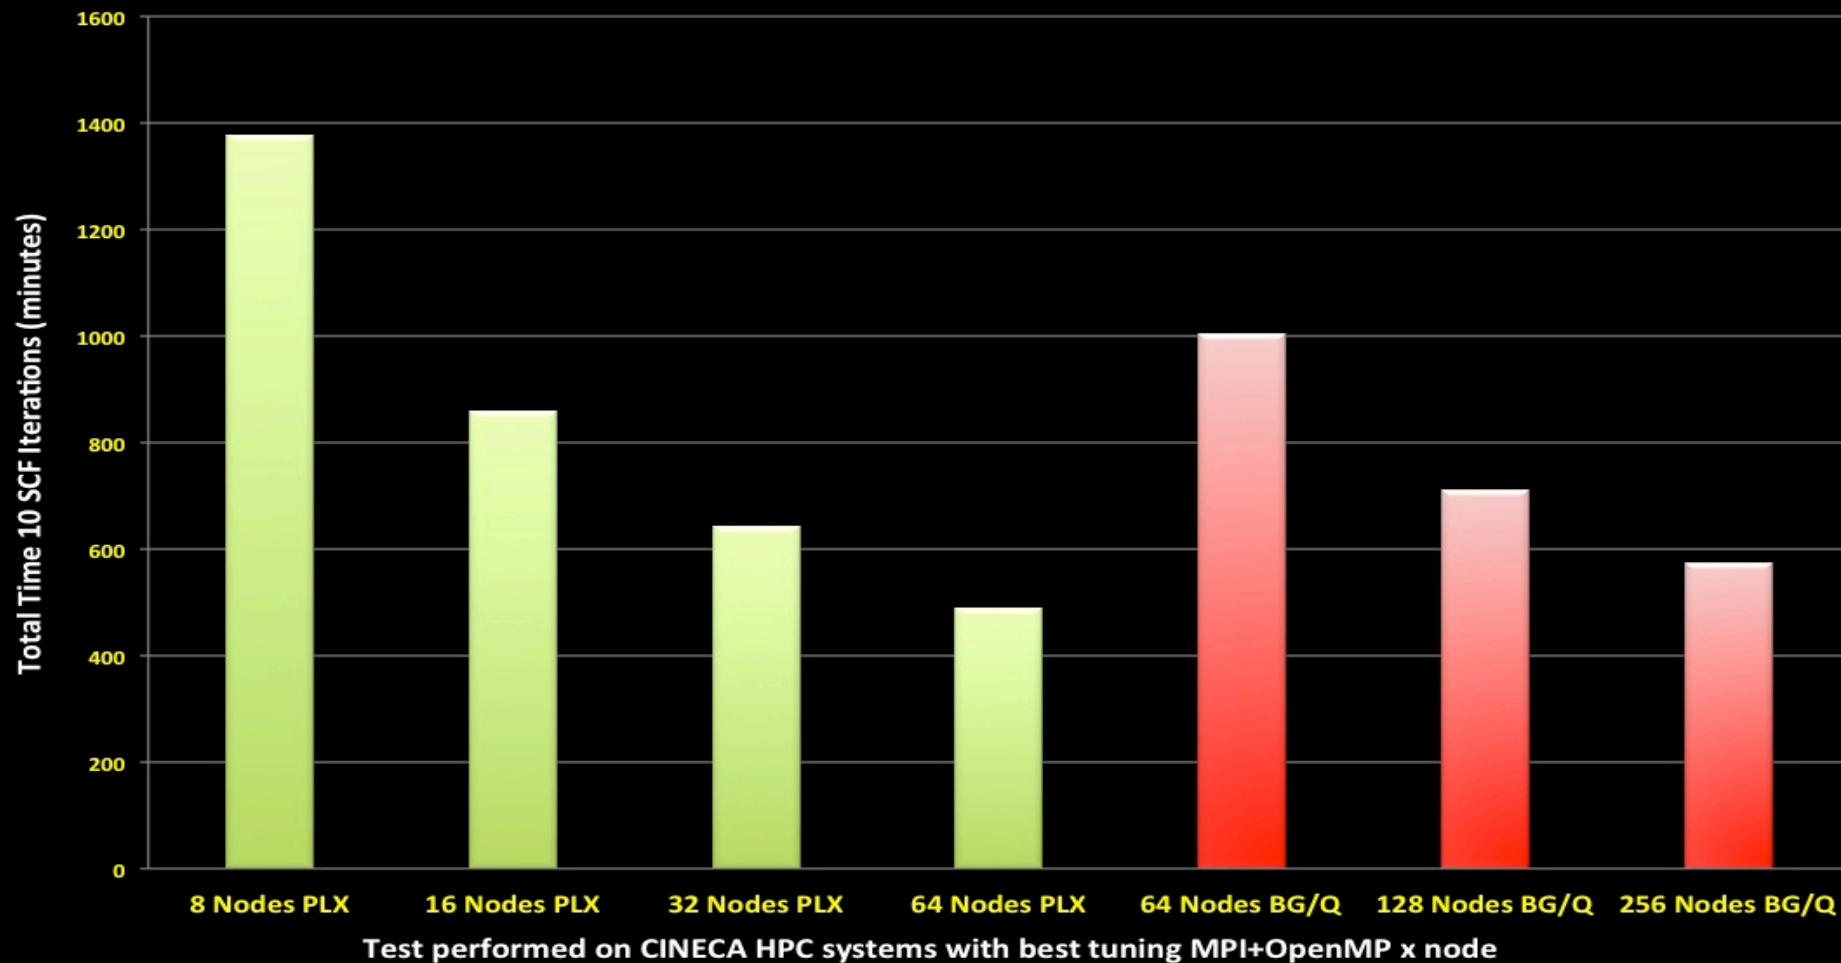
# Best Practice /2

- Scientific case:  
Fe<sub>2</sub>O<sub>3</sub> (120 atoms)
- PI: Dr. Manh Thuong Nguyen,  
Post Doctoral Fellow, The  
Abdus Salam (ICTP)
- Description: Hematite surface,  
1200 electrons, **4 k-points**
- Goal: exploit PWscf on both  
GPU and IBM BG/Q systems





## Hematite surface Fe<sub>2</sub>O<sub>3</sub> (120 Atoms) - 4 K-points





# State of The Art

- If compared with High-End multi-cores platforms, the GPU porting impacts when the communication saturates the MPI traffic and it is inescapable to reduce the number of processes per node
- Better tuning should be done to exploit the High-Throughput model. Huge number of independent systems at the same time
- Few references below:

<https://hpcforge.org/plugins/mediawiki/wiki/pracewp8/images/4/40/MarzariPRACE.pdf>

The high-throughput highway to computational materials design, Stefano Curtarolo, Gus L. W. Hart, Marco Buongiorno Nardelli, Natalio Mingo, Stefano Sanvito & Ohad Levy, Nature Materials 12, 191–201 (2013) doi:10.1038/nmat3568

# Conclusion

- The porting of legacy code is not impossible. But it is a considerable effort.
- USE the GPU PWscf where NVIDIA GPUs are available!!
- The phase of tuning shouldn't scare, it is needed on all the High-End systems!
- Multithreading is inescapable to best exploit the CPU platform

# Further development

- Code enabling on NVIDIA GPUs of current generation
- Feasibility study for improvement of QE-GPU version for other codes the (CP, PH) and EXX section of the PWscf code
- Porting on other accelerated platforms (i.e., Intel-MIC?)



## Acknowledgements:

- Filippo Spiga (Cambridge University/QE-Foundation)
- Paolo Giannozzi (Udine University)
- Carlo Cavazzoni (CINECA)
- Layla Martin-Samos (University of Nova Gorica)
- Rodrigo Neumann Barros Ferreira (Rio de Janeiro Federal University)
- Manh Thuong Nguyen (ICTP)
- Mike Atambo (ICTP)



# References

- P. Giannozzi and et al. QUANTUM ESPRESSO: a modular and open-source software project for quantum simulations of materials. Journal of Physics: Condensed Matter, 21(39), 2009.
- F. Spiga and I. Girotto, "phiGEMM: a CPU-GPU library for porting Quantum ESPRESSO on hybrid systems", Proceeding of 20th Euromicro International Conference on Parallel, Distributed and Network-Based Computing (PDP2012), Special Session on GPU Computing and Hybrid Computing, IEEE Computer Society, (ISBN 978-0-7695-4633-9), pp. 368-375 (2012)
- M. Fatica, "Accelerating LINPACK with CUDA on heterogeneous clusters." GPGPU-2: Proceedings of 2nd Workshop on General Purpose Processing on Graphics Processing Units (New York, NY, USA), ACM, 2009, pp. 46--51.
- Rob Farber, CUDA Application Design and Development, Morgan Kaufmann; 1 edition (November 14, 2011), ISBN-13: 978-0123884268